

# CREACION PLATFORM DRIVER GPIO SWITCH

June 2, 2017

## Part I

# Source Code

```
/*
 * Based on Xilinx GPIO driver
 * Copyright 2008 Xilinx, Inc.
 *
 * This program is free software; you can redistribute
 * it and/or modify
 * it under the terms of the GNU General Public
 * License version 2
 * as published by the Free Software Foundation.
 *
 * You should have received a copy of the GNU General
 * Public License
 * along with this program; if not, write to the Free
 * Software
 * Foundation, Inc., 59 Temple Place, Suite 330,
 * Boston, MA 02111-1307 USA
 */
#include <linux/init.h>
#include <linux/errno.h>
#include <linux/of_device.h>
#include <linux/of_platform.h>
#include <linux/of_gpio.h>
#include <linux/io.h>
#include <linux/gpio.h>
#include <linux/module.h>
#include <linux/slab.h>

struct gcq_gpio_switch_chip {
    struct of_mm_gpio_chip mmchip;
    spinlock_t gpio_lock;    /* Lock used for
                             * synchronization */
};

static int gcq_gpio_switch_get(struct gpio_chip *gc,
    unsigned int gpio) {
    int bank;
    int g;
    struct of_mm_gpio_chip *mm_gc =
        to_of_mm_gpio_chip(gc);

    bank = gpio >> 3;
```

```

        g = gpio%8;
        printk(KERN_INFO "gcq_gpio_get\n");
        return (ioread8(mm_gc->regs + bank) >> g) & 1;
    }

    /*
        status = chip->get_direction(chip, offset);
        if (status > 0) {
            //GPIOF_DIR_IN, or other positive
            status = 1;
            clear_bit(FLAG_IS_OUT, &desc->flags);
        }
        if (status == 0) {
            // GPIOF_DIR_OUT
            set_bit(FLAG_IS_OUT, &desc->flags);
        }
    */
    static int gcq_gpio_switch_get_direction(struct
        gpio_chip *gc, unsigned int gpio) {
        printk(KERN_INFO "gcq_gpio_get_dir\n");
        return 1;
    }

    static int gcq_gpio_switch_probe(struct
        platform_device *pdev) {
        //nodo del dts que se conecta al driver
        struct device_node *node = pdev->dev.of_node;
        int status;
        const u32 *tree_info;
        struct gcq_gpio_switch_chip *gcq_gc;

        gcq_gc = devm_kzalloc(&pdev->dev, sizeof(*
            gcq_gc), GFP_KERNEL);
        if (!gcq_gc)
            return -ENOMEM;

        //inicializamos el lock
        spin_lock_init(&gcq_gc->gpio_lock);
        //valor por defecto de 8 GPIO
        gcq_gc->mmchip.gc.ngpio = 8;
        //leemos el campo del dts xlnx,gpio-width
        tree_info = of_get_property(node, "xlnx,gpio-
            width", NULL);
        if (tree_info)
            gcq_gc->mmchip.gc.ngpio = *tree_info;
    }

```

```

        //creamos puntero a nuestras funciones, mas
        adelante veremos como las llama el kernel
gcq_gc->mmchip.gc.get_direction =
    gcq_gpio_switch_get_direction;
gcq_gc->mmchip.gc.get = gcq_gpio_switch_get;

//registramos el dispositivo GPIO
/* Call the OF gpio helper to setup and
   register the GPIO device */
status = of_mm_gpiochip_add(node, &gcq_gc->
    mmchip);
if (status) {
    dev_err(&pdev->dev, "Failed adding
        memory mapped gpiochip\n");
    return status;
}
platform_set_drvdata(pdev, gcq_gc);
return status;
}

static int gcq_gpio_switch_remove(struct
    platform_device *pdev) {
    struct gcq_gpio_switch_chip *gcq_gc =
        platform_get_drvdata(pdev);
    of_mm_gpiochip_remove(&gcq_gc->mmchip);
    return 0;
}

static struct of_device_id gcq_gpio_switch_of_match[]
    = {
    { .compatible = "dte,gcq-gpio-switch" },
    { /* end of list */ },
};

MODULE_DEVICE_TABLE(of, gcq_gpio_switch_of_match);

static struct platform_driver gcq_gpio_switch_driver =
{
    .driver = {
        .name = "gcq-gpio-switch",
        .of_match_table = of_match_ptr(
            gcq_gpio_switch_of_match),
    },
    .probe = gcq_gpio_switch_probe,
    .remove = gcq_gpio_switch_remove,
};

```

```

static int __init gcq_gpio_switch_init(void) {
    return platform_driver_register(&
        gcq_gpio_switch_driver);
}
subsys_initcall(gcq_gpio_switch_init);

static void __exit gcq_gpio_switch_exit(void) {
    platform_driver_unregister(&
        gcq_gpio_switch_driver);
}
module_exit(gcq_gpio_switch_exit);

MODULE_AUTHOR("GCQ");
MODULE_DESCRIPTION("Testing GPIO SWITCH driver");
MODULE_LICENSE("GPL");

```

## Part II

# DTS

```
switch0: gpio@B8000000 {  
    compatible = "dte,gcq-gpio-switch";  
    reg = <0xB8000000 0x1>;  
    #gpio-cells = <2>;  
    gpio-controller;  
    xlnx,gpio-width = <8>;  
};
```

## Part III

# Registro Driver

```
/*
 * /drivers/base/platform_device.h
 */
#define platform_driver_register(drv)
    __platform_driver_register(drv, THIS_MODULE)

/*
 * /include/linux/export.h
 */
#ifdef MODULE
extern struct module __this_module;
#define THIS_MODULE (&__this_module)
#else
#define THIS_MODULE ((struct module *)0)
#endif

/*
 * /drivers/base/platform.c
 */
/**
 * __platform_driver_register - register a driver for
 * platform-level devices
 * @drv: platform driver structure
 * @owner: owning module/driver
 */
int __platform_driver_register(struct platform_driver
    *drv, struct module *owner) {
    drv->driver.owner = owner;
    drv->driver.bus = &platform_bus_type;
    drv->driver.probe = platform_drv_probe;
    drv->driver.remove = platform_drv_remove;
    drv->driver.shutdown = platform_drv_shutdown;
    return driver_register(&drv->driver);
}

/*
 * /drivers/base/driver.c
 */
/*
 * .driver = {
 *     .name = "gcq-gpio-switch",

```

```

        .of_match_table = of_match_ptr(
            gcq_gpio_switch_of_match),
        -----
        .bus = {
            .name = "
                platform",
            .dev_groups =
                platform_dev_groups
            ,
            .match =
                platform_match, //
                bind platform
                device to platform
                driver.
            .uevent =
                platform_uevent,
            .pm = &
                platform_dev_pm_ops
            ,
        }
    }

    */
int driver_register(struct device_driver *drv) {
    [...]
    //se mira si existe ya un driver con el mismo
    nombre en el platform bus
    other = driver_find(drv->name, drv->bus);
    if (other) {
        printk(KERN_ERR "Error: Driver '%s' is
            already registered, "
            "aborting...\n", drv->name);
        return -EBUSY;
    }
    /*
        principalmente en bus_add_driver se
        llama a:
        1. module_add_driver(drv->owner, drv)
        2. driver_create_file(drv, &
            driver_attr_uevent)
    */
    ret = bus_add_driver(drv);
    if (ret)
        return ret;

    [...]

```



```

        return ret;
    }

void module_add_driver(struct module *mod, struct
    device_driver *drv){
    struct module_kobject *mk = NULL;
    [...]
    mk = &mod->mkobj;
    /**
     *      sysfs_create_link - create
     *      symlink between two objects.
     *      @kobj:  object whose directory
     *              we're creating the link in.
     *      @target:      object we're
     *                    pointing to.
     *      @name:          name of the
     *                    symlink.
     */
    no_warn = sysfs_create_link(&drv->p->kobj, &mk
        ->kobj, "module");
    /*
     * driver_name = "%s:%s", drv->bus->name
     * , drv->name
     */
    driver_name = make_driver_name(drv);
    if (driver_name) {
        /*
         * mk->drivers_dir =
         * kobject_create_and_add("
         * drivers", &mk->kobj);
         */
        module_create_drivers_dir(mk);
        no_warn = sysfs_create_link(mk->
            drivers_dir, &drv->p->kobj,
            driver_name);
        kfree(driver_name);
    }
}

}

/**
 * driver_create_file - create sysfs file for driver.
 * @drv: driver.
 * @attr: driver attribute descriptor.
 */
int driver_create_file(struct device_driver *drv,

```

```

                                const struct driver_attribute *
                                attr) {
    int error;
    if (drv)
        error = sysfs_create_file(&drv->p->
                                kobj, &attr->attr);
    else
        error = -EINVAL;
    return error;
}
EXPORT_SYMBOL_GPL(driver_create_file);

```

## Part IV

# deteccion de GPIO

```
/**
 * of_mm_gpiochip_add - Add memory mapped GPIO chip (
 *                      bank)
 * @np:                device node of the GPIO chip
 * @mm_gc:             pointer to the of_mm_gpio_chip
 *                      allocated structure
 *
 * To use this function you should allocate and fill
 * mm_gc with:
 *
 * 1) In the gpio_chip structure:
 *    - all the callbacks
 *    - of_gpio_n_cells
 *    - of_xlate callback (optional)
 *
 * 3) In the of_mm_gpio_chip structure:
 *    - save_regs callback (optional)
 *
 * If succeeded, this function will map bank's memory
 * and will
 * do all necessary work for you. Then you'll able to
 * use .regs
 * to manage GPIOs from the callbacks.
 */
int of_mm_gpiochip_add(struct device_node *np,
                      struct of_mm_gpio_chip *mm_gc)
{
    int ret = -ENOMEM;
    struct gpio_chip *gc = &mm_gc->gc;
    /**
     * kstrdup - allocate space for and
     *           copy an existing string
     * @s: the string to duplicate
     * @gfp: the GFP mask used in the
     *        kmalloc() call when allocating
     *        memory
     */
    gc->label = kstrdup(np->full_name, GFP_KERNEL);
    ;
    if (!gc->label)
        goto err0;
}
```

```

        //mapeamos la memoria ver mas adelante.
mm_gc->regs = of_iomap(np, 0);
if (!mm_gc->regs)
    goto err1;
gc->base = -1;
//se puede definir la funcion save_regs en el
//driver para dar valor inicial a los
//registros
if (mm_gc->save_regs)
    mm_gc->save_regs(mm_gc);
mm_gc->gc.of_node = np;
ret = gpiochip_add(gc);
if (ret)
    goto err2;
return 0;
err2:
    iounmap(mm_gc->regs);
err1:
    kfree(gc->label);
err0:
    pr_err("%s: GPIO chip registration failed with
        status %d\n",
        np->full_name, ret);
    return ret;
}
EXPORT_SYMBOL(of_mm_gpiochip_add);

```

## mapeado memoria

```
void __iomem *of_iomap(struct device_node *np, int
index) {
    struct resource res;
    if (of_address_to_resource(np, index, &res))
        return NULL;
    //reserva desde (res.start) hasta (res->end -
    res->start + 1)
    return ioremap(res.start, resource_size(&res))
    ;
}
EXPORT_SYMBOL(of_iomap);

int of_address_to_resource(struct device_node *dev,
int index,
                        struct resource *r) {
    const __be32    *addrp;
    u64            size;
    unsigned int    flags;
    const char      *name = NULL;
    addrp = of_get_address(dev, index, &size, &
flags);
    if (addrp == NULL)
        return -EINVAL;
    /* Get optional "reg-names" property to add a
    name to a resource */
    of_property_read_string_index(dev, "reg-names"
, index, &name);
    return __of_address_to_resource(dev, addrp,
size, flags, name, r);
}
EXPORT_SYMBOL_GPL(of_address_to_resource);

const __be32 *of_get_address(struct device_node *dev,
int index, u64 *size,
                        unsigned int *flags) {
    const __be32 *prop;
    unsigned int psize;
    struct device_node *parent;
    struct of_bus *bus;
    int onesize, i, na, ns;
    /* Get parent & match bus type */
    /* devuelve el nodo padre en dts
    parent = of_get_parent(dev);
```

```

if (parent == NULL)
    return NULL;
/*
    *en este punto bus sera
    {
        .name = "default",
        .addresses = "reg",
        .match = NULL,
        .count_cells =
            of_bus_default_count_cells,
        .map = of_bus_default_map,
        .translate = of_bus_default_translate,
        .get_flags = of_bus_default_get_flags,
    }
*/
bus = of_match_bus(parent);
//na = #address-cells
//ns = #size-cells
bus->count_cells(dev, &na, &ns);
of_node_put(parent);
if (!OF_CHECK_ADDR_COUNT(na))
    return NULL;
/* Get "reg" or "assigned-addresses" property
   */
// prop = <0xB8000000 0x1>
prop = of_get_property(dev, bus->addresses, &
    psize);
if (prop == NULL)
    return NULL;
// psize son 2 elementos de u32, 8 bytes
// psize = 8/4 = 2
psize /= 4;
//onesize = 2
onesize = na + ns;
for (i = 0; psize >= onesize; psize -= onesize
    , prop += onesize, i++)
    //index en nuestro caso es 0
    if (i == index) {
        if (size)
            //espacio de memoria
            //necesario le a
            //partir de prop+na,
            //donde comienza ns
            *size = of_read_number
                (prop + na, ns);
        if (flags)

```

```

        *flags = bus->
            get_flags(prop);
        //devolvemos <0xB8000000 0x1>
        return prop;
    }
    return NULL;
}
EXPORT_SYMBOL(of_get_address);

static int __of_address_to_resource(struct device_node
    *dev,
        const __be32 *addrp, u64 size,
        unsigned int flags,
        const char *name, struct resource *r)
    {
    u64 taddr;
    if ((flags & (IORESOURCE_IO | IORESOURCE_MEM))
        == 0)
        return -EINVAL;
    //llama a __of_translate_address con rprop = "
        ranges"
    //en este caso taddr = 0xB8000000
    taddr = of_translate_address(dev, addrp);
    if (taddr == OF_BAD_ADDR)
        return -EINVAL;
    memset(r, 0, sizeof(struct resource));
    if (flags & IORESOURCE_IO) {
        unsigned long port;
        port = pci_address_to_pio(taddr);
        if (port == (unsigned long)-1)
            return -EINVAL;
        r->start = port;
        r->end = port + size - 1;
    } else {
        r->start = taddr;
        r->end = taddr + size - 1;
    }
    r->flags = flags;
    r->name = name ? name : dev->full_name;
    return 0;
}

static u64 __of_translate_address(struct device_node *
    dev,
        const __be32 *

```

```

                                in_addr, const
                                char *rprop) {

struct device_node *parent = NULL;
struct of_bus *bus, *pbus;
__be32 addr[OF_MAX_ADDR_CELLS];
int na, ns, pna, pns;
u64 result = OF_BAD_ADDR;

pr_debug("OF: ** translation for device %s **\n", of_node_full_name(dev));
/* Increase refcount at current level */
of_node_get(dev);
/* Get parent & match bus type */
parent = of_get_parent(dev);
if (parent == NULL)
    goto bail;
bus = of_match_bus(parent);
/* Count address cells & copy address locally
 */
bus->count_cells(dev, &na, &ns);
if (!OF_CHECK_COUNTS(na, ns)) {
    pr_debug("OF: Bad cell count for %s\n", of_node_full_name(dev));
    goto bail;
}
//copia en addr la direccion in_addr = <0
//xB8000000 0x1>
memcpy(addr, in_addr, na * 4);
pr_debug("OF: bus is %s (na=%d, ns=%d) on %s\n",
        bus->name, na, ns, of_node_full_name(
        parent));
of_dump_addr("OF: translating address:", addr,
na);
/* Translate */
for (;;) {
    /* Switch to parent bus */
    of_node_put(dev);
    dev = parent;
    parent = of_get_parent(dev);
    /* If root, we have finished */
    if (parent == NULL) {
        pr_debug("OF: reached root\n");
        //lee el valor desde addr na
        bytes
    }
}

```



```

        result = of_read_number(addr,
                                na);
        break;
    }
    /* Get new parent bus and counts */
    pbus = of_match_bus(parent);
    pbus->count_cells(dev, &pna, &pns);
    if (!OF_CHECK_COUNTS(pna, pns)) {
        printk(KERN_ERR "prom_parse: Bad cell count for %s\n",
               of_node_full_name(dev));
        ;
        break;
    }
    pr_debug("OF: parent bus is %s (na=%d, ns=%d) on %s\n",
            pbus->name, pna, pns,
            of_node_full_name(parent));
    /* Apply bus translation */
    /*en nuestro caso no tenemos la propiedad "rprop" por lo que devuelve 1
    if (of_translate_one(dev, bus, pbus,
                        addr, na, ns, pna, rprop))
        break;
    /* Complete the move up one level */
    na = pna;
    ns = pns;
    bus = pbus;
    of_dump_addr("OF: one level translation:", addr, na);
}
bail:
    of_node_put(parent);
    of_node_put(dev);

    return result; }

```

---

```

static struct of_bus of_busses[] = {
#ifdef CONFIG_OF_ADDRESS_PCI
    /* PCI */
    {
        .name = "pci",
        .addresses = "assigned-addresses",
        .match = of_bus_pci_match,
    }

```

```

        .count_cells = of_bus_pci_count_cells,
        .map = of_bus_pci_map,
        .translate = of_bus_pci_translate,
        .get_flags = of_bus_pci_get_flags,
    },
#endif /* CONFIG_OF_ADDRESS_PCI */
    /* ISA */
    {
        .name = "isa",
        .addresses = "reg",
        .match = of_bus_isa_match,
        .count_cells = of_bus_isa_count_cells,
        .map = of_bus_isa_map,
        .translate = of_bus_isa_translate,
        .get_flags = of_bus_isa_get_flags,
    },
    /* Default */
    {
        .name = "default",
        .addresses = "reg",
        .match = NULL,
        .count_cells =
            of_bus_default_count_cells,
        .map = of_bus_default_map,
        .translate = of_bus_default_translate,
        .get_flags = of_bus_default_get_flags,
    },
};

static struct of_bus *of_match_bus(struct device_node
*np) {
    int i;
    for (i = 0; i < ARRAY_SIZE(of_busses); i++)
        if (!of_busses[i].match || of_busses[i]
            .match(np))
            return &of_busses[i];

    BUG();
    return NULL;
}

/* * Default translator (generic bus) */
static void of_bus_default_count_cells(struct
device_node *dev,
                                     int *addrc, int
                                     *sizec) {

```

```

if (addrc)
    *addrc = of_n_addr_cells(dev); //
    valor de la propiedad #address-
    cells del dts busca desde el hijo
    hasta el nodo raiz
if (sizec)
    *sizec = of_n_size_cells(dev); //
    valor de la propiedad #size-cells
    del dts
}

```

## Registro en sysfs

```
/**
 * gpiochip_add() - register a gpio_chip
 * @chip: the chip to register, with chip->base
 *        initialized
 * Context: potentially before irqs will work
 *
 * Returns a negative errno if the chip can't be
 * registered, such as
 * because the chip->base is invalid or already
 * associated with a
 * different chip. Otherwise it returns zero as a
 * success code.
 *
 * When gpiochip_add() is called very early during
 * boot, so that GPIOs
 * can be freely used, the chip->dev device must be
 * registered before
 * the gpio framework's arch_initcall(). Otherwise
 * sysfs initialization
 * for GPIOs will fail rudely.
 *
 * If chip->base is negative, this requests dynamic
 * assignment of
 * a range of valid GPIOs.
 */
int gpiochip_add(struct gpio_chip *chip) {
    unsigned long    flags;
    int              status = 0;
    unsigned         id;
    //la base es -1
    int              base = chip->base;
    struct gpio_desc *descs;
    //inicializa memoria para un array de chip->
    //ngpio elementos
    //cada uno de tamaño struct gpio_desc
    //inicializa con 0 en sus contenidos
    descs = kcalloc(chip->ngpio, sizeof(descs[0]),
                    GFP_KERNEL);
    if (!descs)
        return -ENOMEM;
    spin_lock_irqsave(&gpio_lock, flags);
    if (base < 0) {
        base = gpiochip_find_base(chip->ngpio)
```

```

        ;
        if (base < 0) {
            status = base;
            spin_unlock_irqrestore(&
                gpio_lock, flags);
            goto err_free_descs;
        }
        chip->base = base;
    }
    status = gpiochip_add_to_list(chip);
    if (status) {
        spin_unlock_irqrestore(&gpio_lock,
            flags);
        goto err_free_descs;
    }
    //valor inicial para la direccion de cada pin
    del GPIO
    /*
    * en la version del kernel v4.11.3
    *
    for (i = 0; i < chip->ngpio; i++) {
        struct gpio_desc *desc = &gdev->descs[
            i];
        desc->gdev = gdev;
        if (chip->get_direction) {
            int dir = chip->get_direction(
                chip, i);
            if (!dir)
                set_bit(FLAG_IS_OUT, &
                    desc->flags);
        } else if (!chip->direction_input) {
            set_bit(FLAG_IS_OUT, &desc->
                flags);
        }
    }
    */
    for (id = 0; id < chip->ngpio; id++) {
        struct gpio_desc *desc = &descs[id];
        desc->chip = chip;
        /* REVISIT: most hardware initializes
        GPIOs as inputs (often
        * with pullups enabled) so power
        usage is minimized. Linux
        * code should set the gpio direction
        first thing; but until
        * it does, and in case chip->

```

```

        get_direction is not set, we may
        * expose the wrong direction in sysfs
        */
        desc->flags = !chip->direction_input ?
            (1 << FLAG_IS_OUT) : 0;
    }
    chip->desc = desc;
    spin_unlock_irqrestore(&gpio_lock, flags);

#ifdef CONFIG_PINCTRL
    INIT_LIST_HEAD(&chip->pin_ranges);
#endif

    if (!chip->owner && chip->dev && chip->dev->
        driver)
        chip->owner = chip->dev->driver->owner
            ;
    status = gpiochip_set_desc_names(chip);
    if (status)
        goto err_remove_from_list;
    status = of_gpiochip_add(chip);
    if (status)
        goto err_remove_chip;
    acpi_gpiochip_add(chip);
    status = gpiochip_sysfs_register(chip);
    if (status)
        goto err_remove_chip;
    pr_debug("%s: registered GPIOs %d to %d on %s\n", __func__,
        chip->base, chip->base + chip->ngpio -
            1,
        chip->label ? : "generic");
    return 0;

err_remove_chip:
    acpi_gpiochip_remove(chip);
    gpiochip_free_hogs(chip);
    of_gpiochip_remove(chip);
err_remove_from_list:
    spin_lock_irqsave(&gpio_lock, flags);
    list_del(&chip->list);
    spin_unlock_irqrestore(&gpio_lock, flags);
    chip->desc = NULL;
err_free_descs:
    kfree(descs);

```

```

        /* failures here can mean systems won't boot
        ... */
pr_err("%s: GPIOs %d..%d (%s) failed to
    register\n", __func__,
        chip->base, chip->base + chip->ngpio -
            1,
        chip->label ? : "generic");
return status;
}
EXPORT_SYMBOL_GPL(gpiochip_add);

```

---

```

static int gpiochip_find_base(int ngpio) {
    struct gpio_chip *chip;
    //ARCH_NR_GPIO = 512 por defecto
    // base inicial = 512-8 = 504
    int base = ARCH_NR_GPIO - ngpio;
    /**
        * list_for_each_entry_reverse -
            iterate backwards over list of
            given type.
        * @pos:      the type * to use as a
            loop cursor. // chip struct
            gpio_chip
        * @head:      the head for your list
            . //gpio_chips list_head
        * @member:     the name of the
            list_head within the struct.
        recorre los existentes GPIO, por lo
        tanto en GPIO0 no se itera y su
        base es 504.
        veamos para GPIO1
    */
    list_for_each_entry_reverse(chip, &gpio_chips,
        list) {
        /* found a free space? para GPIO1*/
        //chip es GPIO0, 504+8 <= 504, es
        falso
        if (chip->base + chip->ngpio <= base)
            break;
        else
            /* nope, check the space right
            before the chip */
            //GPIO 1 base = baseGPIO0 - 8

```

```

        = 496
        base = chip->base - ngpio;
    }
    if (gpio_is_valid(base)) {
        pr_debug("%s: found new base at %d\n",
            __func__, base);
        return base;
    } else {
        pr_err("%s: cannot find free range\n",
            __func__);
        return -ENOSPC;
    }
}

```

---

```

/*
 * Add a new chip to the global chips list, keeping
 * the list of chips sorted
 * by base order.
 *
 * Return -EBUSY if the new chip overlaps with some
 * other chip's integer
 * space.
 */
static int gpiochip_add_to_list(struct gpio_chip *chip
    ) {
    struct list_head *pos;
    struct gpio_chip *_chip;
    int err = 0;
    /* find where to insert our chip */
    //gpio_chips es una list_head, linked list con
    // todos los GPIO del sistema.
    //veremos el ejemplo con GPIO1
    //gpio_chips = [&gpio_chips, GPIO0, &gpio_chips]
    list_for_each(pos, &gpio_chips) {
        // _chip es GPIO0
        _chip = list_entry(pos, struct
            gpio_chip, list);
        /* shall we insert before _chip? */
        //GPIO0 base = 504
        //GPIO1 base = 496
        //504 >= 496+8 es cierto
        if (_chip->base >= chip->base + chip->
            ngpio)

```



```

        break;
    }
    /* are we stepping on the chip right before?
       */
    //pos != &gpio_chips es si hemos llegado al
       final
    //pos->prev != &gpio_chips indica que la lista
       no es vacia
    //si next y prev == &gpio_chips significa que
       la lista es vacia
    //pos = GPIO0, pos->prev = &gpio_chips
    if (pos != &gpio_chips && pos->prev != &
        gpio_chips) {
        _chip = list_entry(pos->prev, struct
            gpio_chip, list);
        if (_chip->base + _chip->ngpio > chip
            ->base) {
            dev_err(chip->dev,
                "GPIO integer space
                overlap, cannot add
                chip\n");
            err = -EBUSY;
        }
    }
    if (!err)
        //gpio_chips = [&gpio_chips, GPIO1,
            GPIO0, &gpio_chips]
        list_add_tail(&chip->list, pos);
    return err;
}

```

---

```

//registramos el GPIO en sysfs
int gpiochip_sysfs_register(struct gpio_chip *chip) {
    struct device *dev;
    /*
    * Many systems add gpio chips for SOC support
    very early,
    * before driver model support is available.
    In those cases we
    * register later, in gpiolib_sysfs_init() ...
    here we just
    * verify that _some_ field of gpio_class got
    initialized.

```

```

    */
    if (!gpio_class.p)
        return 0;
    /* use chip->base for the ID; it's already
       known to be unique */
    dev = device_create_with_groups(&gpio_class,
        chip->dev, MKDEV(0, 0),
                                chip,
                                gpiochip_groups,
                                ,
                                "gpiochip%d",
                                chip->base)
                                ;

    if (IS_ERR(dev))
        return PTR_ERR(dev);
    mutex_lock(&sysfs_lock);
    chip->cdev = dev;
    mutex_unlock(&sysfs_lock);
    return 0;
}

```

## Part V

# System Calls mediante sysfs

## UTILIDADES

```
/*
 * /sys/class/gpio/gpiochipN/
 * /base ... matching gpio_chip.base (N)
 * /label ... matching gpio_chip.label
 * /ngpio ... matching gpio_chip.ngpio
 */
static ssize_t base_show(struct device *dev,
                        struct device_attribute
                        *attr, char *buf) {
    const struct gpio_chip *chip =
        dev_get_drvdata(dev);
    return sprintf(buf, "%d\n", chip->base);
}
static DEVICE_ATTR_RO(base);

static ssize_t label_show(struct device *dev,
                        struct device_attribute
                        *attr, char *buf) {
    const struct gpio_chip *chip =
        dev_get_drvdata(dev);
    return sprintf(buf, "%s\n", chip->label ? : ""
    );
}
static DEVICE_ATTR_RO(label);

static ssize_t ngpio_show(struct device *dev,
                        struct device_attribute
                        *attr, char *buf) {
    const struct gpio_chip *chip =
        dev_get_drvdata(dev);
    return sprintf(buf, "%u\n", chip->ngpio);
}
static DEVICE_ATTR_RO(ngpio);
```

## EXPORT

```
/*
 * /sys/class/gpio/export ... write-only
 *     integer N ... number of GPIO to export (full
 *         access)
 * /sys/class/gpio/unexport ... write-only
 *     integer N ... number of GPIO to unexport
 */
static ssize_t export_store(struct class *class,
                           struct class_attribute
                               *attr,
                           const char *buf,
                           size_t len) {
    long
    struct gpio_desc
    int
        gpio;
        *desc;
        status;

    //string to int
    status = kstrtoul(buf, 0, &gpio);
    if (status < 0)
        goto done;

    desc = gpio_to_desc(gpio);
    /* reject invalid GPIOs */
    if (!desc) {
        pr_warn("%s: invalid GPIO %ld\n",
                __func__, gpio);
        return -EINVAL;
    }
    /* No extra locking here; FLAG_SYSFS just
     * signifies that the
     * request and export were done by on behalf
     * of userspace, so
     * they may be undone on its behalf too.
     */
    status = gpiod_request(desc, "sysfs");
    if (status < 0) {
        if (status == -EPROBE_DEFER)
            status = -ENODEV;
        goto done;
    }
    status = gpiod_export(desc, true);
    if (status < 0)
        gpiod_free(desc);
}
```

```

        else
            set_bit(FLAG_SYSFS, &desc->flags);

done:
    if (status)
        pr_debug("%s: status %d\n", __func__,
                status);
    return status ? : len;
}

-----

struct gpio_desc *gpio_to_desc(unsigned gpio) {
    struct gpio_chip *chip;
    unsigned long flags;

    spin_lock_irqsave(&gpio_lock, flags);
    //GPIO0 pin 3 = 504 + 3 = 507
    list_for_each_entry(chip, &gpio_chips, list) {
        //GPIO0 base = 504 gpio = 507 ngpio =
        8
        if (chip->base <= gpio && chip->base +
            chip->ngpio > gpio) {
            spin_unlock_irqrestore(&
                gpio_lock, flags);
            return &chip->desc[gpio - chip
                ->base];
        }
    }
    spin_unlock_irqrestore(&gpio_lock, flags);
    if (!gpio_is_valid(gpio))
        WARN(1, "invalid GPIO %d\n", gpio);
    return NULL;
}

int gpiod_export(struct gpio_desc *desc, bool
    direction_may_change){
    struct gpio_chip      *chip;
    struct gpiod_data      *data;
    [...]
    chip = desc->chip;
    [...]
    data = kzalloc(sizeof(*data), GFP_KERNEL);
    data->desc = desc;
    mutex_init(&data->mutex);
    // si en nuestro driver hemos creado funciones

```

```

        para direction_input y direction_output
        // estamos estableciendo que el GPIO puede
        cambiar su direccion y ser I/O
        if (chip->direction_input && chip->
            direction_output)
            data->direction_can_change =
                direction_may_change;
        else
            data->direction_can_change = false;

        [...]
        //creamos el gpio en /sys con gpio%d
        descriptor, GPIO0 3 = gpio507
        dev = device_create_with_groups(&gpio_class,
            chip->dev,
                                MKDEV(0, 0),
                                data,
                                gpio_groups
                                ,
                                ioname ?
                                ioname : "
                                gpio%u",
                                desc_to_gpio(
                                desc));
    }

```

## DIRECCION

```
int gpiod_get_direction(struct gpio_desc *desc) {
    struct gpio_chip    *chip;
    unsigned             offset;
    int                  status = -EINVAL;

    chip = gpiod_to_chip(desc);
    offset = gpio_chip_hwgpio(desc);
    if (!chip->get_direction)
        return status;

    //vemos la definicion de get_direction
    //implementada por nuestro driver
    status = chip->get_direction(chip, offset);
    if (status > 0) {
        /* GPIOF_DIR_IN, or other positive */
        status = 1;
        clear_bit(FLAG_IS_OUT, &desc->flags);
    }
    if (status == 0) {
        /* GPIOF_DIR_OUT */
        set_bit(FLAG_IS_OUT, &desc->flags);
    }
    return status;
}

static ssize_t direction_show(struct device *dev,
                             struct device_attribute *attr, char *
                             buf) {
    struct gpiod_data *data = dev_get_drvdata(dev);
    ;
    struct gpio_desc *desc = data->desc;
    ssize_t             status;

    mutex_lock(&data->mutex);
    gpiod_get_direction(desc);
    status = sprintf(buf, "%s\n",
                    test_bit(FLAG_IS_OUT, &desc->
                    flags)? "out" : "in");
    mutex_unlock(&data->mutex);
    return status;
}
```

```

-----

/* * gpiod_direction_input - set the GPIO direction to
    input */
int gpiod_direction_input(struct gpio_desc *desc) {
    struct gpio_chip      *chip;
    int                    status = -EINVAL;

    if (!desc || !desc->chip) {
        pr_warn("%s: invalid GPIO\n", __func__);
        return -EINVAL;
    }

    chip = desc->chip;
    // se debe haber implementado en el driver el
    // getValue y el direction_input
    if (!chip->get || !chip->direction_input) {
        gpiod_warn(desc,
            "%s: missing get() or direction_input() operations\n",
            __func__);
        return -EIO;
    }
    //hacemos la preparacion en nuestro driver
    status = chip->direction_input(chip,
        gpio_chip_hwgpio(desc));
    if (status == 0)
        clear_bit(FLAG_IS_OUT, &desc->flags);
        //se cambia el valor en el GPIO

    trace_gpio_direction(desc_to_gpio(desc), 1,
        status);
    return status;
}

static int _gpiod_direction_output_raw(struct
    gpio_desc *desc, int value) {
    struct gpio_chip      *chip;
    int                    status = -EINVAL;

    /* GPIOs used for IRQs shall not be set as
        output */
    if (test_bit(FLAG_USED_AS_IRQ, &desc->flags))
        {

```



```

        gpiod_err(desc,
            "%s: tried to set a GPIO tied to an IRQ as output\n",
            __func__);
        return -EIO;
    }
    /* Open drain pin should not be driven to 1 */
    if (value && test_bit(FLAG_OPEN_DRAIN, &desc->flags))
        return gpiod_direction_input(desc);
    /* Open source pin should not be driven to 0 */
    if (!value && test_bit(FLAG_OPEN_SOURCE, &desc->flags))
        return gpiod_direction_input(desc);

    chip = desc->chip;
    /* se debe haber implementado en el driver el setValue y el direction_output */
    if (!chip->set || !chip->direction_output) {
        gpiod_warn(desc,
            "%s: missing set() or direction_output() operations\n",
            __func__);
        return -EIO;
    }
    status = chip->direction_output(chip,
        gpio_chip_hwgpio(desc), value);
    if (status == 0)
        set_bit(FLAG_IS_OUT, &desc->flags);
    trace_gpio_value(desc_to_gpio(desc), 0, value);
    ;
    trace_gpio_direction(desc_to_gpio(desc), 0, status);
    return status;
}

//se cambia la direccion y se inicializa su valor
static ssize_t direction_store(struct device *dev,
    struct device_attribute *attr, const
    char *buf, size_t size) {
    struct gpiod_data *data = dev_get_drvdata(dev);
    ;

```

```

struct gpio_desc *desc = data->desc;
ssize_t          status;

mutex_lock(&data->mutex);
if (sysfs_streq(buf, "high"))
    status = gpiod_direction_output_raw(
        desc, 1);
else if (sysfs_streq(buf, "out") ||
        sysfs_streq(buf, "low"))
    status = gpiod_direction_output_raw(
        desc, 0);
else if (sysfs_streq(buf, "in"))
    status = gpiod_direction_input(desc);
else
    status = -EINVAL;
mutex_unlock(&data->mutex);
return status ? : size;
}

```

## VALOR

```
//llamada del sysfs
static ssize_t value_show(struct device *dev,
                          struct device_attribute *attr, char *
                          buf) {
    struct gpiod_data *data = dev_get_drvdata(dev);
    ;
    struct gpio_desc *desc = data->desc;
    ssize_t status;

    mutex_lock(&data->mutex);
    status = sprintf(buf, "%d\n",
                    gpiod_get_value_cansleep(desc));
    mutex_unlock(&data->mutex);
    return status;
}

int gpiod_get_value_cansleep(const struct gpio_desc *
desc) {
    int value;

    might_sleep_if(extra_checks);
    if (!desc)
        return 0;

    value = _gpiod_get_raw_value(desc);
    if (value < 0)
        return value;
    if (test_bit(FLAG_ACTIVE_LOW, &desc->flags))
        value = !value;
    return value;
}

static int _gpiod_get_raw_value(const struct gpio_desc
*desc) {
    struct gpio_chip *chip;
    int offset;    int value;
    chip = desc->chip;
    offset = gpio_chip_hwgpio(desc);
    //llamamos a la funcion get de nuestro driver
    value = chip->get ? chip->get(chip, offset) :
        -EIO;
    /*
     * FIXME: fix all drivers to clamp to [0,1] or
```

```

        return negative,
        * then change this to:
        * value = value < 0 ? value : !!value;
        * so we can properly propagate error codes.
        */
value = !!value;
trace_gpio_value(desc_to_gpio(desc), 1, value)
;
return value;
}

```

-----

```

static ssize_t value_store(struct device *dev,
                          struct device_attribute *attr, const
                          char *buf, size_t size) {
    struct gpiod_data *data = dev_get_drvdata(dev)
    ;
    struct gpio_desc *desc = data->desc;
    ssize_t                status;

    mutex_lock(&data->mutex);
    //comprobar que el GPIO sea de salida
    if (!test_bit(FLAG_IS_OUT, &desc->flags)) {
        status = -EPERM;
    } else {
        long                value;
        status = kstrtol(buf, 0, &value);
        if (status == 0) {
            gpiod_set_value_cansleep(desc,
                                     value);
            status = size;
        }
    }
    mutex_unlock(&data->mutex);
    return status;
}

```

```

void gpiod_set_value_cansleep(struct gpio_desc *desc,
int value) {
    might_sleep_if(extra_checks);
    if (!desc)
        return;
    if (test_bit(FLAG_ACTIVE_LOW, &desc->flags))
        value = !value;
}

```

```

        _gpiod_set_raw_value(desc, value);
    }

static void _gpiod_set_raw_value(struct gpio_desc *
    desc, bool value) {
    struct gpio_chip      *chip;
    chip = desc->chip;
    trace_gpio_value(desc_to_gpio(desc), 0, value)
        ;

    if (test_bit(FLAG_OPEN_DRAIN, &desc->flags))
        _gpio_set_open_drain_value(desc, value
            );
    else if (test_bit(FLAG_OPEN_SOURCE, &desc->
        flags))
        _gpio_set_open_source_value(desc,
            value);
    else
        chip->set(chip, gpio_chip_hwgpio(desc)
            , value); //por defecto llama a la
            funcion set de nuestro driver
}

```